

AI Media Studio

User Manual

A fully local, AI-powered pipeline that turns one topic into a finished, narrated, subtitled short-form video. No cloud services and no paid APIs.

Version	0.1.0
Requires	Python 3.12 or newer
Runs on	Windows, verified on Windows 11
Output	1080x1920 H.264 MP4 with AAC audio

Contents

- 1 What this does
- 2 Before you start
- 3 Installing
- 4 Configuration
- 5 Generating a video
- 6 Resuming a run
- 7 Inside a project folder
- 8 Scenes, motion and transitions
- 9 Animating scenes with Stable Video Diffusion
- 10 How timing works
- 11 Background music
- 12 Setting up the ComfyUI workflow
- 13 Troubleshooting
- 14 Running the tests

1 What this does

You give AI Media Studio a topic. It researches and writes a script, breaks it into ordered scenes, narrates them, generates an image for each scene, writes subtitles, and renders the whole thing into a vertical MP4 ready for Reels, Shorts or TikTok. Everything runs on your own machine.

The pipeline

Topic	
-> Script and ordered scenes	(Ollama)
-> Narration, measured word by word	(Kokoro)
-> Scene durations reconciled to speech	
-> One image per scene	(ComfyUI)
-> One clip per scene, optional	(ComfyUI, section 9)
-> Subtitles	(SRT)
-> Final MP4	(FFmpeg)

Each stage writes its results to a project folder as it goes. Nothing is held only in memory, which is what makes an interrupted run resumable (section 6).

What you need to supply

- A topic, as a short phrase.
- Optionally a style, a voice, background music, and whether you want subtitles.

2 Before you start

Four local services do the actual work. Install and start these first.

Component	Purpose	Notes
Ollama	Writes the script and plans scenes	Must be running. Pull a model first, for example <code>ollama pull llama3.1:8b</code> .
ComfyUI	Generates one image per scene, and optionally animates it	Must be running, with an image model installed. See section 12, and section 9 for animation.
Kokoro	Speaks the narration	Installed automatically as a Python dependency. Downloads its voice on first use.
FFmpeg	Renders the final video	Either on your PATH or given as a full path in the configuration.

Hardware

Image generation is the demanding part. A machine with an 8 GB NVIDIA GPU produces an SDXL image in roughly 20 to 35 seconds once the model is loaded in memory, and 90 seconds or more before that. Less video memory still works but is slower, because the model is reloaded for each image.

3 Installing

```
pip install -e .
copy config\settings.example.toml config\settings.toml
```

Then open `config\settings.toml` and replace the two placeholder values: the Ollama model name and the Kokoro voice. Nothing will run until you do, because the placeholders are not real model names.

Installing without -e? A non-editable install has no project folder to read from, so tell it where your settings file is with `--config` or the `AI_MEDIA_CONFIG` environment variable, and use absolute paths inside that file.

Checking it worked

```
ai-media-studio generate --help
```

4 Configuration

All settings live in one TOML file. By default that is `config\settings.toml` inside the project folder.

Choosing a different file

Highest priority first:

- The `--config` option on the command line.
- The `AI_MEDIA_CONFIG` environment variable.
- The project's own `config\settings.toml`.

The settings that matter most

Setting	What it controls	Default
<code>ollama.model</code>	Which local language model writes the script	none, you must set it
<code>kokoro.voice</code>	Which narration voice is used	none, you must set it
<code>kokoro.speed</code>	Speaking rate	1.0
<code>kokoro.scene_tail_padding_seconds</code>	Silence after each scene so speech does not butt against a cut	0.2
<code>paths.ffmpeg_executable</code>	FFmpeg command or full path	ffmpeg
<code>paths.output_dir</code>	Where project folders are created	output
<code>comfyui.workflow_path</code>	The ComfyUI workflow JSON to run	config/comfyui_workflow.json
<code>comfyui.timeout_seconds</code>	How long to wait for one image	120
<code>comfyui.max_retries</code>	Retries after a recoverable image failure	2
<code>video.width, video.height</code>	Output resolution	1080 by 1920
<code>video.frames_per_second</code>	Output frame rate	30
<code>video.transition_duration_seconds</code>	How long a transition lasts	0.5
<code>video.animate_still_scenes</code>	Move the image even when the storyboard asks for no camera motion	true

Setting	What it controls	Default
video.camera_motion_strength	How far a zoom or pan travels	0.2
video.render_timeout_seconds	How long to wait for the render	300
video.target_duration_seconds	Aim the finished video at this many seconds. 0 lets the script run to any length	0
video.clip_smoothing	How the frames a slowed clip lacks are made. Only used with section 9	blend
svd.enabled	Animate each scene image instead of panning across it	false
svd.frames, svd.fps	Length of a generated clip, as frames over frames per second	25 over 6
svd.motion_bucket_id	How much a clip moves, at the cost of coherence	127
subtitles.max_characters_per_cue	Longest caption shown at once	32
gpu.device	Device for models running in this process, meaning narration	auto
music.volume	Background music level	0.15
music.ducking_ratio	How hard music dips under narration	6.0
logging.level	Detail written to the log	INFO

Raise `comfyui.timeout_seconds` **to about 300**. The very first image of a run also has to load the image model into memory, which alone can take longer than the 120 second default. Otherwise the first scene times out and is retried needlessly.

Overriding settings without editing the file

Any of these environment variables takes precedence over the file, which is convenient for one-off changes and for scripting.

AI_MEDIA_CONFIG	AI_MEDIA_LOG_LEVEL
AI_MEDIA_OLLAMA_BASE_URL	AI_MEDIA_LOG_DIRECTORY
AI_MEDIA_OLLAMA_MODEL	AI_MEDIA_LOG_CONSOLE_ENABLED
AI_MEDIA_OLLAMA_TIMEOUT_SECONDS	AI_MEDIA_LOG_FILE_ENABLED
AI_MEDIA_COMFYUI_BASE_URL	AI_MEDIA_GPU_DEVICE
AI_MEDIA_COMFYUI_WORKFLOW_PATH	AI_MEDIA_VIDEO_WIDTH
AI_MEDIA_COMFYUI_TIMEOUT_SECONDS	AI_MEDIA_VIDEO_HEIGHT
AI_MEDIA_KOKORO_VOICE	AI_MEDIA_VIDEO_FPS
AI_MEDIA_KOKORO_SPEED	AI_MEDIA_VIDEO_RENDER_TIMEOUT_SECONDS
AI_MEDIA_VIDEO_TRANSITION_DURATION_SECONDS	

Settings without a variable in this list, including everything under `[svd]` and `[subtitles]`, are set in the file.

5 Generating a video

```
ai-media-studio generate --topic "Why Japan Never Sleeps"
```

With every option in use. The trailing backtick continues a command across lines in PowerShell; in Command Prompt use a caret instead.

```
ai-media-studio generate `
  --topic "Why Japan Never Sleeps" `
  --style "cinematic documentary" `
  --voice "af_heart" `
  --output "D:\Media Projects" `
  --music `
  --subtitle
```

Options

Option	Meaning
--topic	The subject of the video. Required unless you are resuming.
--resume	Continue an existing project folder instead of starting fresh. Cannot be combined with --topic.
--style	Narrative and visual style, for example <i>cinematic documentary</i> . Passed to the script writer.
--voice	Override the narration voice for this run only.
--output	Where the project folder is created, overriding the configured output folder.
--config	Use a different settings file.
--music	Pick a random track from your music folder and mix it under the narration.
--subtitle	Write subtitles and burn them into the video.

What you will see

```
[#####-----] 1/3 Generating storyboard, narration, and images
[#####-----] 2/3 Generating subtitles
[#####-----] 3/3 Rendering final MP4

Generation complete
Generation time: 505.94 seconds
Provider versions: ollama=unknown
Final output: C:\AI-Media-Studio\output\20260808T164700953894Z_why-japan-never-sleeps\video\final.mp4
```

Roughly how long it takes

Measured on a laptop with an 8 GB NVIDIA GPU, for a six scene video.

Stage	Typical time	Notes
Script	15 to 60 seconds	Depends on the language model and topic.
Narration	80 to 115 seconds	Slower the first time, while the voice loads.

Stage	Typical time	Notes
Images	20 to 95 seconds each	The first two are slowest, then it speeds up.
Clips	about 4 minutes each	Only when [svd] is enabled. See section 9.
Subtitles	under a second	Written from measured word timings.
Rendering	5 to 25 seconds	Longer with music, subtitles and clips.

A whole run is typically 6 to 10 minutes, and almost all of it is image generation. Enabling animation (section 9) takes it to roughly half an hour. If a run fails near the end, resume it rather than starting over.

6 Resuming a run

Every stage saves its work as it finishes, so a run that stops partway through can carry on. This matters most when the image service becomes unavailable mid-run, which otherwise throws away a finished script and narration.

```
ai-media-studio generate --resume "output\20260808T164700953894Z_why-japan-never-sleeps" --subtitle
```

When resuming:

- The script is read back from the project folder and never regenerated.
- Existing narration is reused as it is, along with the word timings recorded when it was spoken, so captions still follow the narration word by word.
- Scene images already on disk are kept, and only missing ones are generated.
- Animated clips are kept too. They are the most expensive thing a run produces, so they are never made twice.
- The video is always rendered again, which is quick.

Resuming a project that is already complete regenerates nothing and simply re-renders, so it is a cheap way to try different music or subtitle settings without paying for image generation again.

To regenerate one artifact, delete it and resume. Removing `narration.wav` re-speaks the script and re-times every scene to it; removing a single file from `images\` regenerates only that scene.

7 Inside a project folder

Each run creates one timestamped folder under your output folder, named after the time and the topic.

```

20260808T164700953894Z_why-japan-never-sleeps\
manifest.json      what ran, which providers, how long, and whether it succeeded
storyboard.json   the script and every scene, with their final durations
narration.wav     the spoken narration for the whole video
word_timings.json when each word is spoken, which the captions are built from
subtitles.srt     subtitle cues, when --subtitle was used
images\
  scene_001.png   one image per scene
  scene_002.png
clips\
  scene_001.webm  one animated clip per scene, when [svd] is enabled
video\
  final.mp4       the finished video
logs\

```

manifest.json is the place to look when something went wrong. It records which stage failed and why, and is written even for a failed run.

storyboard.json holds the narration text and the scene timings actually used. If you want to know why a video is the length it is, read this.

word_timings.json is written when the narration is spoken, because only the voice can measure it. Deleting it does not break anything, but captions then fall back to one cue per scene, a paragraph at a time.

8 Scenes, motion and transitions

The script writer assigns each scene a camera motion and a transition into the next one. Only the values below are accepted; anything else the model invents is converted to the closest match, so an unusual word never spoils a run.

Camera motion

Value	Effect
none	The image is held still.
zoom_in	Slow push in towards the centre.
zoom_out	Slow pull back from the centre.
pan, pan_right	Drifts across the image to the right.
pan_left	Drifts across the image to the left.

Transitions

Value	Effect
cut, none	An instant change with no blend.
fade, crossfade	One scene blends into the next.
dissolve	A softer, more textured blend.
wipeleft, wiperight	The next scene slides across horizontally.
wipeup, wipedown	The next scene slides across vertically.

Transition length comes from `video.transition_duration_seconds`. A transition is automatically shortened if either neighbouring scene is too brief to accommodate it.

Language models almost always ask for no camera motion, which leaves every image frozen on screen. Scenes asking for none are therefore given a movement anyway, alternating between zoom and pan so consecutive scenes differ. Set `video.animate_still_scenes` to false to follow the storyboard exactly.

The opening line

The script writer names a hook for the opening seconds. It leads the first scene's narration, so it is spoken and captioned like any other line. The merge happens once, before the storyboard is saved, so resuming a run does not speak it twice, and a hook the model already used as its opening line is left alone rather than repeated.

A model asked only for a hook writes atmospheric scene setting, so the script prompt says what the hook is for: one sentence, at most twelve words, stating a fact, question or claim rather than describing the view.

Read the opening line before you publish. The prompt asks for an accurate hook and the model still occasionally overstates one, because nothing checks a claim against reality. Automatic fact checking is not built yet.

Subtitles

Captions follow the narration word by word rather than showing a whole scene's text at once. The voice provider records when each word is spoken while it synthesizes the speech, and those timings drive the cues. Sentences are divided into cues of roughly equal length, so none is left holding a single word, and each stays on screen until the next begins.

Those timings are saved to the project folder, so resuming a run keeps captions word by word rather than dropping back to one cue per scene.

9 Animating scenes with Stable Video Diffusion

By default each scene is a still image with a slow pan or zoom across it. Turn this on and the picture itself moves: rain falls, light flickers, and clouds drift. Each scene is generated twice, once as a still image and again as a short clip animating that image.

This is off by default because it is slow and needs another model. It uses the same ComfyUI you already have.

Setting it up

- Put a Stable Video Diffusion checkpoint in ComfyUI's `models\checkpoints` folder, named by `config\svd_workflow.json`.
- Set `svd.enabled` to true.

```
huggingface-cli download stabilityai/stable-video-diffusion-img2vid-xt \
  svd_xt.safetensors --local-dir comfyui\models\checkpoints
```


What it costs

Expect three to four minutes per scene on an 8 GB card, so a five scene video takes around half an hour rather than eight minutes. Clips are saved in the project folder and reused when a run is resumed, because they are by far the most expensive thing a run produces.

On a card with less memory than the model needs, the model is reloaded for every single clip and the time per scene barely improves as the run goes on. This is normal and not a fault.

Why clips are slowed, and what that needs

A clip runs `svd.frames` divided by `svd.fps` seconds, about four by default, while a scene lasts as long as its narration. Clips are therefore slowed to fill their scene.

Slowing 25 frames across a ten second scene would leave each picture on screen for four tenths of a second, which looks like stuttering rather than slow motion. The frames in between are generated instead, controlled by `video.clip_smoothing`.

Value	Cost per scene	Result
blend	a few seconds	Each frame fades into the next. The default, and hard to tell from <i>motion</i> on the soft drifting movement these clips usually have.
motion	one to two minutes	Movement between frames is followed, which is sharper when a clip has a clearly moving subject. Raise <code>video.render_timeout_seconds</code> to use it.
none	none	Frames are repeated, which stutters. Included for comparison.

Camera motion is not applied to a scene that has a clip, because the picture already moves. A scene whose animation fails falls back to its still image, so a failure costs quality rather than the whole run.

10 How timing works

This is worth understanding, because it explains why your video is the length it is.

The language model estimates how long each scene should take to say. Those estimates are unreliable and are never used to time the video. Instead the narration for each scene is generated first and measured, and those measurements become the timeline.

The result is that:

- The video is exactly as long as the narration. There is no trailing silence, and narration is never cut off mid-sentence.
- Each image is on screen for exactly as long as its own narration.
- Each transition begins the moment its scene stops speaking.
- Subtitles line up with the speech, because they use the same measurements.

Aiming at a length

Because the video runs exactly as long as the narration, a length can only be decided when the script is written. Set `video.target_duration_seconds` and the script writer is asked for a matching number of words, phrased per scene, which a language model follows far better than a total for the whole script.

Asking is not enough on its own, so the result is checked. How long a script takes to speak is known from its word count, which costs nothing to work out, so a script that would run too long or too short is thrown away and asked for again, up to three times, keeping the closest attempt. Nothing is narrated or illustrated until the length fits, so a miss costs one model call rather than a whole run.

Leave it at **0**, the default, and a script runs to whatever length it wants. Thirty seconds or under suits discovery on most platforms.

11 Background music

Put audio files in the folder named by `music.directory`, which is `music\` by default. Accepted formats are mp3, wav, m4a, aac, flac and ogg.

With `--music`, one track is chosen at random, looped to fit, faded in and out, and automatically dipped underneath the narration so the voice stays clear. Adjust the level with `music.volume` and the strength of the dip with `music.ducking_ratio`.

If the folder is missing or empty you get a warning and the video is still produced, just without music.

12 Setting up the ComfyUI workflow

Image generation runs whatever ComfyUI workflow you point it at, so you can use any model and any settings ComfyUI supports. Export your workflow from ComfyUI in API format and set `comfyui.workflow_path` to it.

Your workflow must satisfy two rules:

- Exactly one positive prompt node, found by following the sampler's positive input. The scene's image prompt is written into it for each image.
- Exactly one Save Image node, which is where the finished image is collected.

If either is missing or ambiguous you get a clear configuration error naming the problem, rather than a failed render later on. Nothing else about the workflow matters, so resolution, sampler, steps and model are entirely yours to choose.

Generate portrait images that match your output aspect ratio, for example 768 by 1344 for a 1080 by 1920 video. Images are scaled and cropped to fit, so a landscape image will lose a lot of its edges.

13 Troubleshooting

The log file, by default `logs\ai_media_studio.log`, records every stage with timings and the reason for any failure.

Message you see	What it means and what to do
Configuration file not found	There is no settings file where it looked. Copy the example file to <code>config\settings.toml</code> , or pass <code>--config</code> .
Generation failed during storyboard: HTTP Error 404	Ollama is running but does not have the model named in your settings. Check <code>ollama list</code> and pull the model, or correct <code>ollama.model</code> .

Message you see	What it means and what to do
Generation failed during storyboard: connection_failed	Ollama is not running, or is on a different address than <code>ollama.base_url</code> .
Generation failed during image: Unable to reach ComfyUI	ComfyUI is not running or has stopped. Start it, then resume the run with <code>--resume</code> so the script and narration are not lost.
ComfyUI generation timed out	Usually the first image of a run, while the model loads. Raise <code>comfyui.timeout_seconds</code> to about 300.
No positive prompt node could be discovered	Your workflow has no prompt node reachable from a sampler's positive input, or has several. See section 12.
Multiple Save Image output nodes were discovered	Your workflow saves more than one image. Leave exactly one Save Image node.
Generation failed during render: FFmpeg is unavailable	FFmpeg is not on your PATH. Set <code>paths.ffmpeg_executable</code> to its full path.
Generation failed during render: FFmpeg rendering timed out	Raise <code>video.render_timeout_seconds</code> . Generating the frames between a clip's own frames is the slow part, especially with <code>video.clip_smoothing</code> set to <i>motion</i> .
Generation failed during clip	ComfyUI could not animate a scene. Check that the Stable Video Diffusion checkpoint named in <code>config\svd_workflow.json</code> is actually in ComfyUI's checkpoints folder. Resume the run afterwards; the clips already made are kept.
Narration runs N seconds but the scenes span M seconds	The narration and the scene list no longer agree, which happens if one is regenerated without the other. Delete <code>narration.wav</code> and resume, so both are rebuilt together.
Music warning, or Subtitle warning	Not a failure. The video is still produced without that extra. Usually an empty or missing music folder.

Image generation is unusually slow

If every image takes 90 seconds rather than settling to around 25, the image model is being reloaded for each one because memory is tight. Close other applications, or use a smaller image model. Speed usually improves a few scenes into a run.

The picture jumps rather than moving smoothly

If animated scenes look like they are lagging, check that `video.clip_smoothing` is not set to *none*. A clip holds about twenty five frames, so filling a long scene by repeating them leaves each picture on screen for a noticeable moment. See section 9.

ComfyUI stops on its own during long runs

On machines with limited memory the image service can exit without an error. Restart it and resume the run, which keeps everything already generated.

14 Running the tests

The test suite uses only the Python standard library and never contacts a local model, so it runs in well under a second and is safe to run at any time.

```
python -m unittest discover -s tests -t .
```

Use it to confirm an installation is sound before spending time on a full run.